

6.046

Lecture 23

Parallel Algorithms

①

Multithreaded algorithmsMatrix Multiplication ($n \times n$) $C = AB \rightarrow$ divide and conquer (not Strassen)

$$\begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix} = \begin{pmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{pmatrix} + \begin{pmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{22}B_{21} & A_{22}B_{22} \end{pmatrix}$$

Mult (C, A, B, n) // n is power of 2temp matrix $T[1 \dots n, 1 \dots n]$ if $n=1$ then $C[1,1] \leftarrow A[1,1] \cdot B[1,1]$ else \leftarrow partition matrices // $O(1)$ timespawn Mult($C_{11}, A_{11}, B_{11}, n/2$)spawn Mult($C_{12}, A_{11}, B_{12}, n/2$)spawn Mult($C_{21}, A_{21}, B_{11}, n/2$)spawn Mult($C_{22}, A_{21}, B_{12}, n/2$)spawn Mult($T_{11}, A_{12}, B_{21}, n/2$)spawn Mult($T_{12}, A_{12}, B_{22}, n/2$)spawn Mult($T_{21}, A_{22}, B_{21}, n/2$)spawn Mult($T_{22}, A_{22}, B_{22}, n/2$)syncAdd(C, T, n)

$$\begin{pmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{pmatrix}$$

$$\begin{pmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{22}B_{21} & A_{22}B_{22} \end{pmatrix}$$

Add (C, T, n) // $C \leftarrow C + T$

< base & partitioning >

spawn Add ($C_{11}, T_{11}, n/2$)

spawn Add ($C_{12}, T_{12}, n/2$)

spawn Add ($C_{21}, T_{21}, n/2$)

spawn Add ($C_{22}, T_{22}, n/2$)

sync

Analysis

Let $M_P(n)$ = P-processor execution time for Mult

$A_P(n)$ = P-processor execution time for Add

Work (want to run preprogram on 1 processor in same time as ~~non-parallel~~ ~~parallel~~ programs)

$$A_1(n) = 4 A_1\left(\frac{n}{2}\right) + \Theta(1) = \Theta(n^2) \quad \text{1st case Master}$$

↑
good, *
input matrix is n^2 size
same as serial

$$M_1(n) = 8 M_1\left(\frac{n}{2}\right) + \Theta(n^2) \\ = \Theta(n^3) \quad \text{1st case Master}$$

↖ same as serial program (not 8 spawns)

Critical path length

$$A_\infty(n) = A_\infty\left(\frac{n}{2}\right) + \Theta(1) \\ = \Theta(\lg n) \quad \text{case 2, Master}$$

$$M_\infty(n) = M_\infty\left(\frac{n}{2}\right) + \Theta(\lg n)$$

$$= \Theta(\lg^2 n) \quad \text{case 2, Master}$$

↖ not max with all spawns

because it depends on their execution

(sync)

↪ adding across sync

when analysing

work: + across
spawns

critical
path: max across
spawns

6.046

Lecture 23

Parallelism

$$\bar{P} = \frac{M_1(n)}{M_\infty(n)} = \Theta\left(\frac{n^3}{\lg^2 n}\right) \leftarrow \begin{array}{l} \text{linear} \\ \text{expect speed up if} \\ \text{run with upto } \bar{P} \text{ processors} \end{array}$$

For 1000×1000 matrices
assume constants irrelevant

$$\bar{P} \approx \frac{1000^3}{10^2} = 10^7 \quad (10 \text{ million processors}) \leftarrow \text{Blue Gene } > 10^6 \text{ processors}$$

\bar{P} is much bigger than typical P

\hookrightarrow expect linear speedup in typical settings

Trade parallelism for space efficiency (constant matter,)
a lot of parallelism

Mult-Add (C, A, B, n) // $C \leftarrow C + A \cdot B$

< base + partition >

spawn Mult-Add ($C_{11}, A_{11}, B_{11}, n/2$)

\leftarrow need to initialize as well in this case

spawn Mult-Add ($C_{22}, A_{21}, B_{12}, n/2$)

sync
spawn Mult-Add ($C_{11}, A_{12}, B_{21}, n/2$)

spawn Mult-Add ($C_{22}, A_{22}, B_{22}, n/2$)

sync

work

$$MA_1(n) = \Theta(n^3)$$

same recurrence with $\Theta(1)$ instead of $\Theta(n^2)$ for add.

Critical path length

\leftarrow add maxes across syncs

$$MA_\infty(n) = 2 MA_\infty\left(\frac{n}{2}\right) + \Theta(1)$$

$$= \Theta(n) \text{ case 1, Master}$$

Parallelism

$$\bar{P} = \frac{MA_1(n)}{MA_2(n)} = \Theta\left(\frac{n^3}{n}\right) = \Theta(n^2)$$

For 1000×1000 matrices $\bar{P} \approx 10^6$ (enough for typical settings)
decreased \bar{P} from $\approx 10^7$ to $\approx 10^6$

and decreased space by eliminating T temp matrix

Faster in practice, since less space.

~~the~~ Approach: given a lot of parallelism
trade it for other aspects,
even constants that make a difference
in practice

Sorting

Merge-Sort (A, p, r) // sort $A[p..r]$

if $p < r$

then $q \leftarrow \lfloor (p+r)/2 \rfloor$

spawn Merge-Sort (A, p, q)

spawn Merge-Sort ($A, q+1, r$)

sync

Merge (A, p, q, r) // merge $A[p..q]$ with
 $A[q+1..r]$

else (base case)

Work: $T_1(n) = 2T_1\left(\frac{n}{2}\right) + \Theta(n) = \Theta(n \lg n)$



↳ same as serial

CPL: $T_\infty(n) = T_\infty\left(\frac{n}{2}\right) + \Theta(n) = \Theta(n)$

Parallelism: $\bar{P} = \frac{T_1(n)}{T_\infty(n)} = \Theta(\lg n)$ Not much parallelism
at all!
Punny!

Must parallelize Merge

2 sorted arrays



① getting $\frac{l}{2}$ from larger array guarantees that recursion is at least $\frac{n}{4} \Rightarrow$ guarantee of better splits

P-Merge ($A[1 \dots l]$, $B[1 \dots m]$, $C[1 \dots n]$)

" Merge A & B into C ; $n = l + m$

< wlog $l > m$ >

< base > " if small could do serial merge in practice

find j such that $B[j] \leq A[l/2] \leq B[j+1]$ using binary search

spawn P-Merge ($A[1 \dots l/2]$, $B[1 \dots j]$, $C[1 \dots l/2 + j]$)

spawn P-Merge ($A[l/2 + 1 \dots l]$, $B[j+1 \dots m]$, $C[l/2 + j + 1 \dots n]$)

sync

Critical path length

$$PM_{\infty}(n) \leq PM_{\infty}\left(\frac{3n}{4}\right) + \Theta(\lg n) = \Theta(\lg^2 n) \quad \text{case 2 Master}$$

Work

$$PM_1(n) = PM_1(\alpha n) + PM_1((1-\alpha)n) + \Theta(\lg n)$$

where $\frac{1}{4} \leq \alpha \leq \frac{3}{4}$

Substitution:

$$PM_1(k) \leq ak - b \lg k, \text{ where } a, b > 0$$

$$PM_1(n) \leq a(\alpha n) - b \lg(\alpha n) + a((1-\alpha)n) - b \lg((1-\alpha)n) + \Theta(\lg n)$$

$$= an - b(\lg(\alpha n) + \lg((1-\alpha)n)) + \Theta(\lg n)$$

$$= an - b(\lg \alpha + \lg n + \lg(1-\alpha) + \lg n) + \Theta(\lg n)$$

← binary search

by IH since $\frac{1}{4} \leq \alpha \leq \frac{3}{4}$

$$\begin{aligned}
 &= an - b(\lg d + \lg n + \lg(1-d) + \lg n) + \Theta(\lg n) \\
 &= an - b\lg n - (b(\lg n + \lg(d(1-d)))) - \Theta(\lg n)
 \end{aligned}$$

$\leq an - b\lg n$ if we choose b large enough such that

$b(\lg n + \lg(d(1-d)))$ dominates $\Theta(\lg n)$

choose a big enough to satisfy base of induction.

Thus $PM_1(n) = \Theta(n) \leftarrow$ only $O(n)$ shown above
 $\mathcal{O}(n)$ is also true

Merge-Sort Analysis (parallelized Merge)

Work: $T_1(n) = \Theta(n \lg n) \leftarrow$ Merge work did not change compared to the non-parallelized Merge

CPL: $T_\infty(n) = T_\infty(\frac{n}{2}) + \Theta(\lg^2 n)$ Merge-Sort recurrence
 $= \Theta(\lg^3 n)$

$\bar{P} = \Theta(\frac{n}{\lg^2 n})$ Best is $\bar{P} = \Theta(\frac{n}{\lg n})$

Problem in practice is to get constant down.