


Modern Memory hierarchy

Random access model: access in memory at  $O(1)$ . Different Model.



- 1) slower ~~access~~ (higher latency)
  - 2) and bigger space
  - 3) bandwidth should be same
- the larger space, the more time to access speed of light as fundamental limit.

Cost to Access

$$= \text{latency} + \frac{\text{amount of data}}{\text{bandwidth}}$$

↑  
limited by speed of light

↖ rate at which can get data out, assume const/fixed

Idea: as latency goes up, increase the amount of data, then the amortized cost to access an element (fixed bandwidth) goes down

$$\frac{\text{Amortized cost to access one element}}{\text{amount}} = \frac{\text{latency}}{\text{amount}} + \frac{1}{\text{bandwidth}}$$

↑ latency amortized by amount / size of block

Spatial locality

want algorithms to use all elements in a block (after getting it from disk into memory, slow to faster)

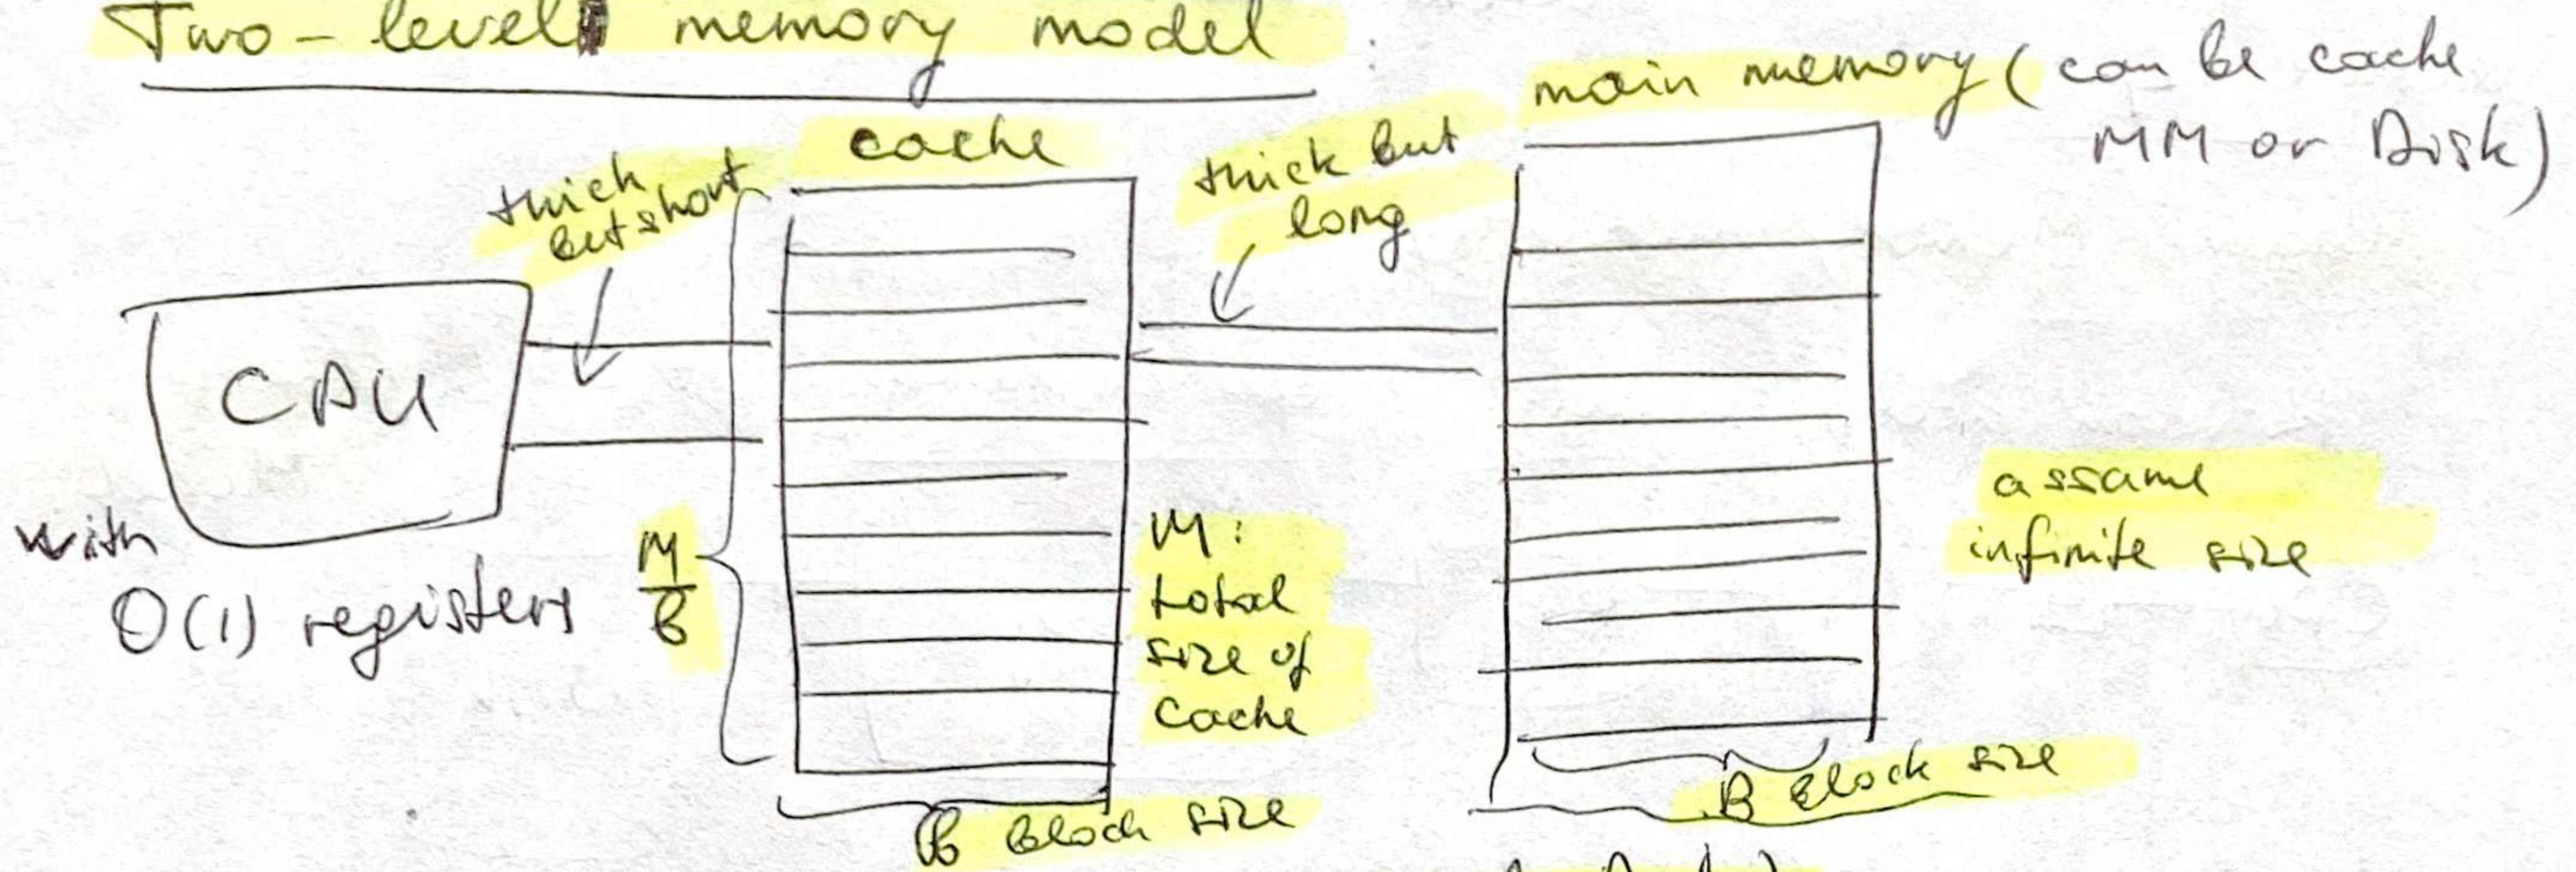
Temporal locality

want ideally to reuse blocks (if algorithm is above linear will use elements in input more than once)




(cache-aware model)

## Two-level memory model



can represent any 2 levels in (incl. Disk)



slower (higher latency)  
more space  
bandwidth same

Assume CPU can access cache instantaneously (free)  
If CPU needs data, check cache if there get it free,  
else get entire block from main memory,  
potentially need to ~~remove~~ remove data from cache  
and write it back to main memory to free  
up cache.

- accesses to cache free (but still measure computation time)
- count block memory transfer between cache and main memory

memory transfer: read or write a block  
from/into main memory

$$MT(N) (= MT_{B,M}(N))$$

$\uparrow$  size of problem  
 $\uparrow$  block size  
 $\uparrow$  total size of cache } fixed parameters

(the only variable that can change)  $\rightarrow$  cannot recurse on  
Block size or total cache size (fixed) parameters



Cache aware algorithm:  $B$  fixed (know fixed cache parameters)

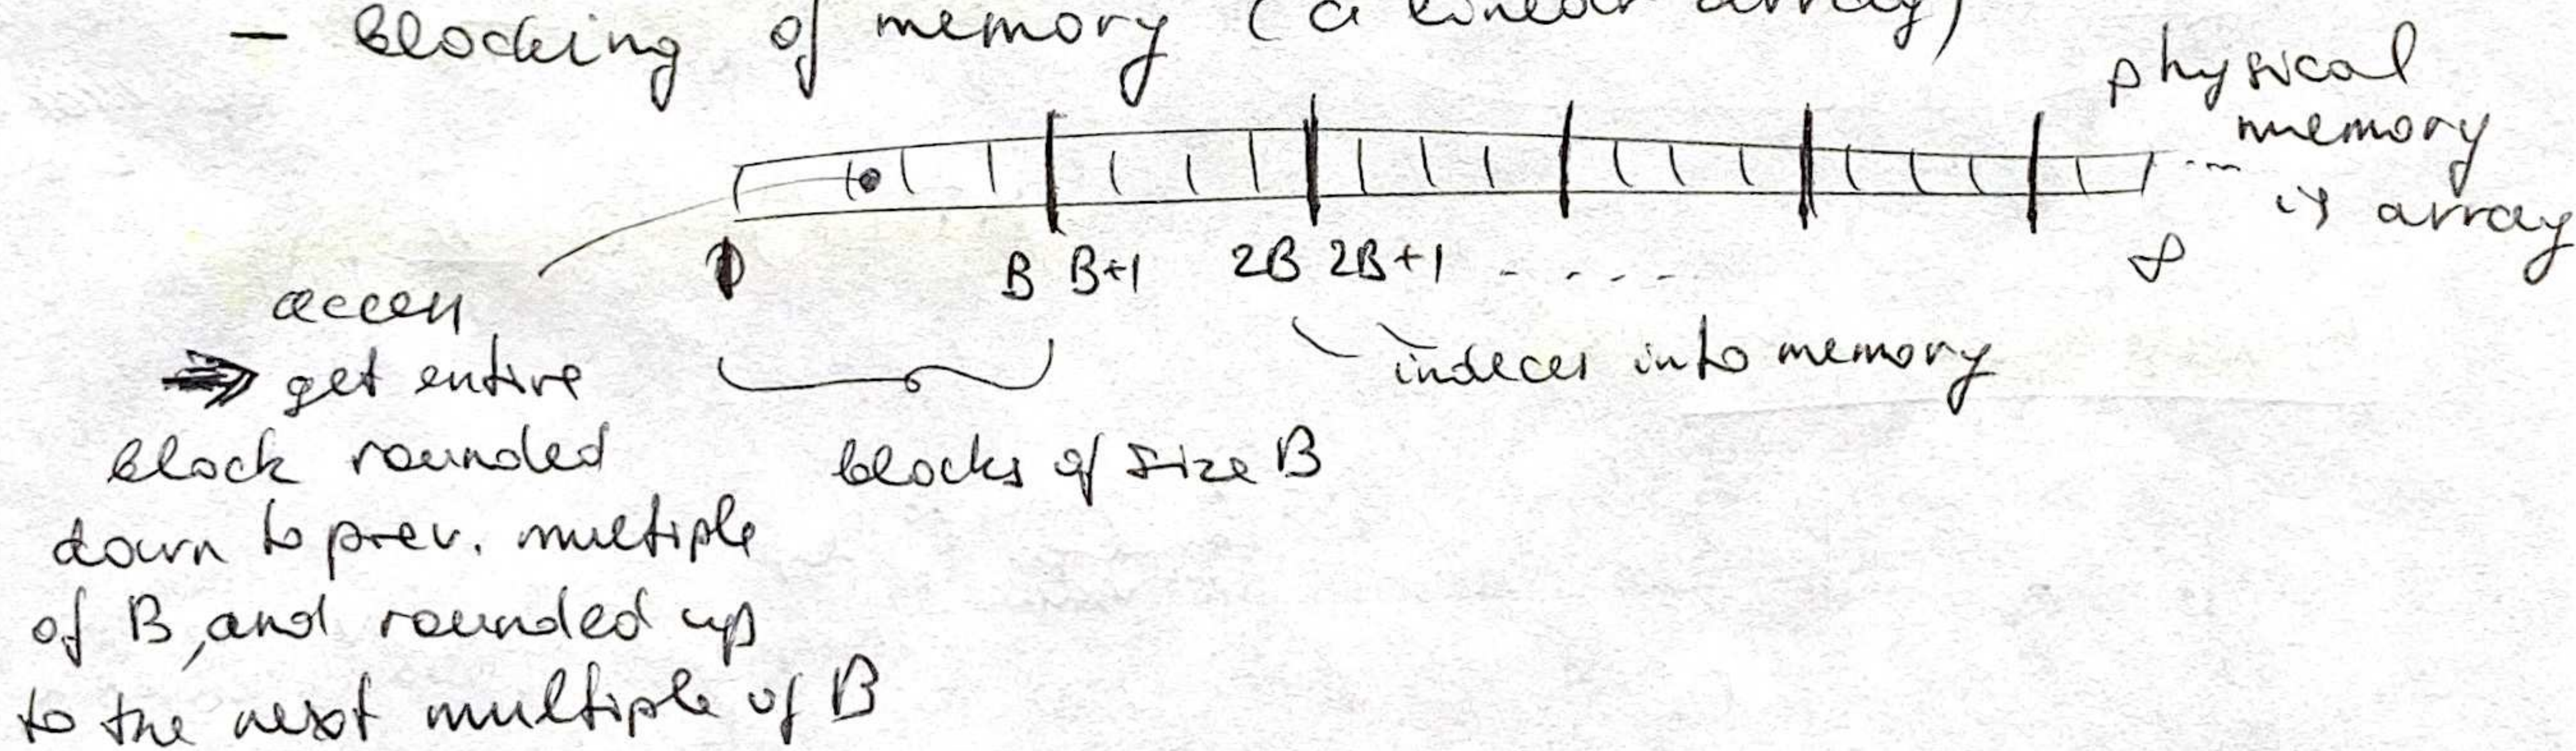
Cache-oblivious: want to design algorithms that perform well no matter the values of  $B$  and  $M$

motivation

- 1) most "just memory" assuming algorithms are already cache-oblivious.
- 2) some will do well in the two-level memory model and some will not.

Cache-oblivious algorithm:

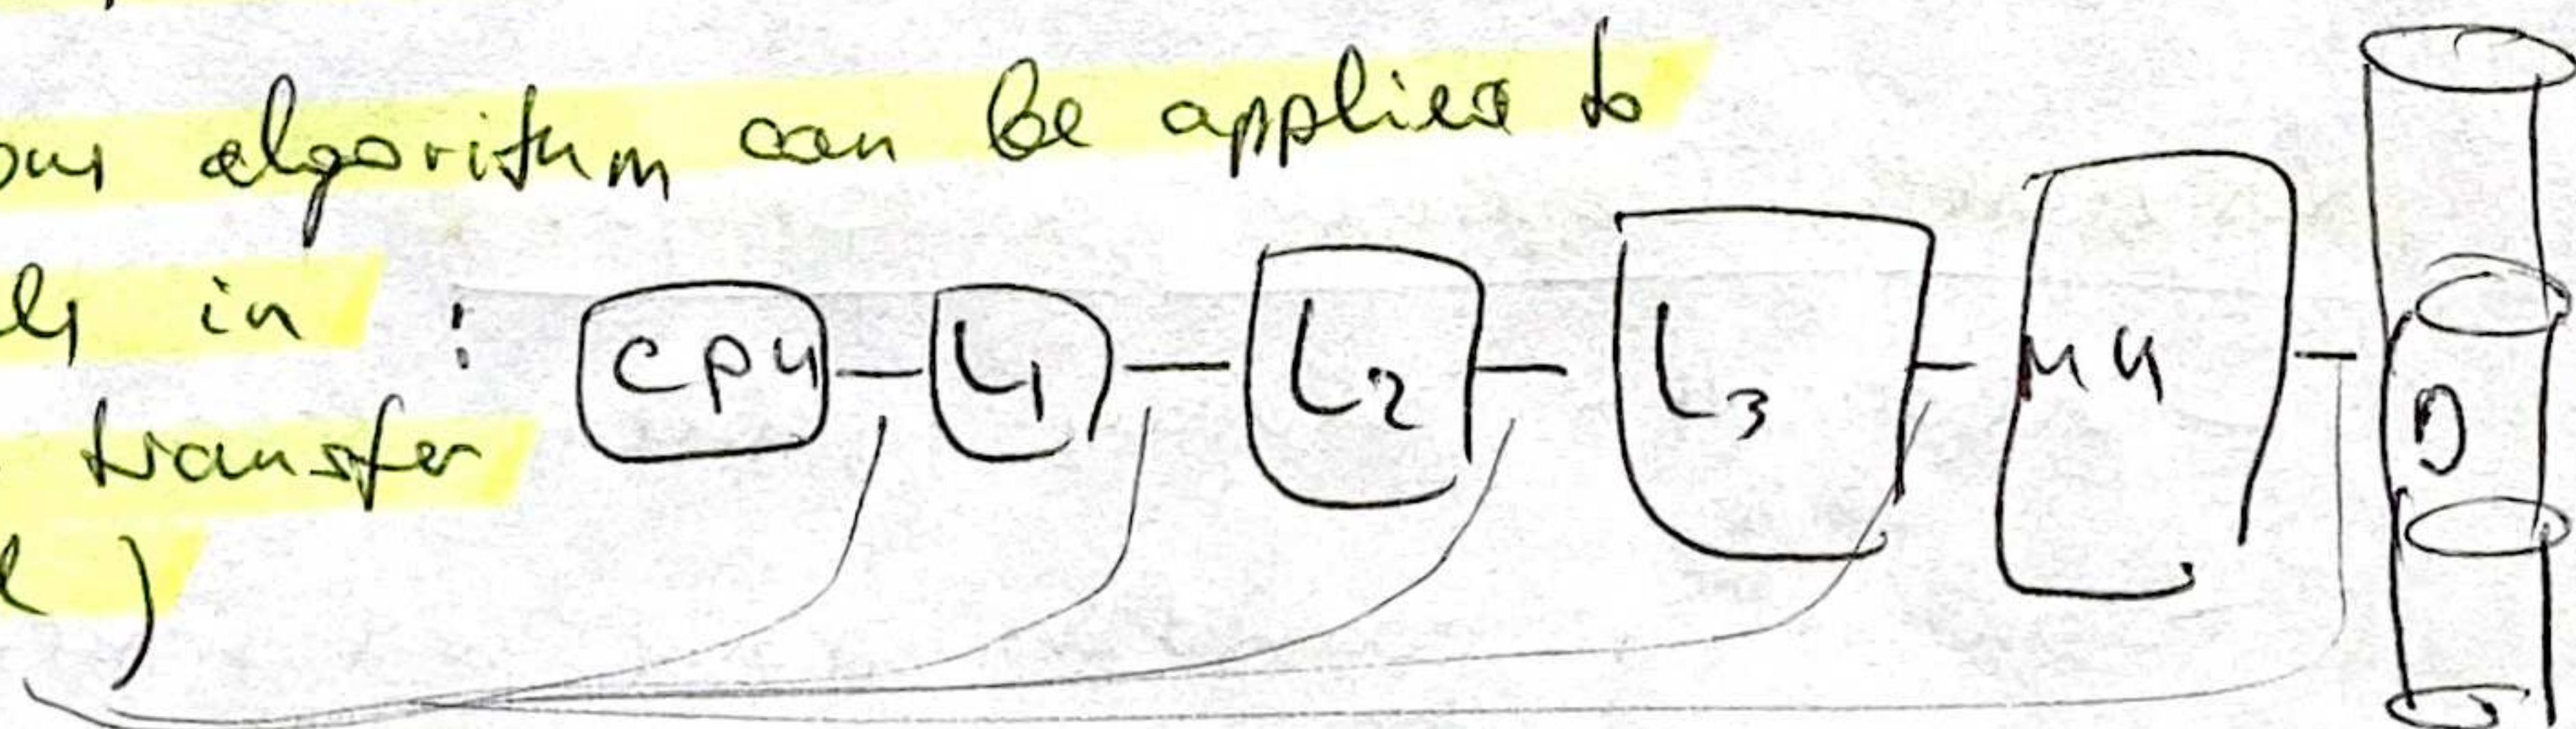
- algorithm does not know  $B$  &  $M$
- accessing element automatically fetch block containing it
- evicts block that will be used furthest in future
- blocking of memory (a linear array)



Fact:

[ if efficient on 2 levels  $\Rightarrow$  efficient on  $l$  levels ]  
 & cache oblivious

Efficient cache oblivious algorithm can be applied to any pair of levels in (i.e. any memory transfer level)




Because cache-oblivious algs do not tune to particular values of  $B$  &  $M$ .



## Basic algorithms

Scanning (A, W) // visiting items in array in order  
e.g. sum array

for  $i \leftarrow 1$  to  $N$   
do visit  $A[i]$




$$MT(N) = O\left(\frac{N}{B} + 1\right)$$

$N$  could be  $< B$ ,  
more precisely  
should be 2 since  
first and last block  
can be non-full

$O(1)$  parallel scans

Reverse (A, N):


for  $i \leftarrow 1$  to  $\lfloor N/2 \rfloor$   
do exchange  $A[i] \leftrightarrow A[N-i+1]$



Assuming  $\frac{M}{B} \geq 2$  (cache can  
store ~~at least~~ at least  
2 blocks)

$$MT(N) = O\left(\frac{N}{B} + 1\right)$$

Binary Search (x)



if  $x$  equals to  
compared  
 $O(1)$  values

hope:  $\log_B N$  without  
knowing  $B$

$$MT(N) = \cancel{O\left(\lg\left(\frac{N}{B} + 1\right)\right)} = O(\lg N - \lg B + 1) \quad \text{bad}$$

once search  
narrows down  
to 1 block

Divide & Conquer algorithms (incl. Binary search)

- algorithm divides problem down to  $O(1)$  size

- analysis considers point at which:


change base  
case of  
recurrence

- problem fits in cache ( $\leq M$ )
- problem fits in  $O(1)$  blocks



Order Statistics (median)

- ① conceptually partition array into  $\frac{N}{5}$  5-tuples
- ② compute median of each tuple
- ③ recursively compute median of these medians  $\times \{MT(\frac{N}{5})\}$
- ④ partition around  $x$
- ⑤ recurse on one side  $MT(\frac{3}{4}N)$

Analysis

$$MT(N) = MT\left(\frac{N}{5}\right) + MT\left(\frac{3}{4}N\right) + O\left(\frac{N}{B} + 1\right)$$

$$MT(1) \leftarrow \text{bad}$$

# leaves:  $L(N) = L\left(\frac{N}{5}\right) + L\left(\frac{3}{4}N\right)$

$$L(1) = 1$$

$$N^{\alpha} = \left(\frac{N}{5}\right)^{\alpha} + \left(\frac{3}{4}N\right)^{\alpha}$$

$$1 = \left(\frac{1}{5}\right)^{\alpha} + \left(\frac{3}{4}\right)^{\alpha} \quad \alpha \approx 0.8398$$

$$\Rightarrow L(N) = \left(\frac{N}{5}\right)^{\alpha} + \left(\frac{3}{4}N\right)^{\alpha} = \omega\left(\frac{N}{B}\right) \quad \text{bad}$$


$$MT(B) = O(1)$$

# leaves  $= \left(\frac{N}{B}\right)^{\alpha} = O\left(\frac{N}{B}\right)$

cost roughly geometric down the tree  
 $\Rightarrow$  root dominates.

$$\Rightarrow MT(N) = O\left(\frac{N}{B}\right)$$


$\nwarrow$  can't do better  
 because need to read  
 all  $N$ .





# Matrix Multiplication:

Standard



$$C = A \cdot B$$

$O(\frac{N}{B})$  mem. transfers to compute  $c_{ij}$

$$\Rightarrow MT(N) = O(\frac{N^3}{B})$$

col. in B moves  
for each value in A's row segment  
but want  $O(\frac{N^2}{B})$

under memory layout assumption we have a  
good spatial locality  
bad temporal locality

Memory layout  
assume C stored in  
row-major

A row-major

B col-major

then for each row in C  
must transfer a row in A  
and entire B

## Block algorithm


$$\begin{matrix} \frac{n}{2} \\ n \\ \frac{n}{2} \end{matrix} \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix}$$

$$\approx \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}$$

$$\begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}$$

$$= \begin{pmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{pmatrix} + \begin{pmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{22}B_{21} & A_{22}B_{22} \end{pmatrix}$$

Store matrices recursively by block in memory.



same layout for  
C and B

recursively in memory  
in one array of size N

← powerful idea in  
cache-oblivious algorithms




$$MT(N) = 8 MT\left(\frac{N}{2}\right) + O\left(\frac{N^2}{B}\right)$$

relies on  
recursive memory  
layout

scan  
same consecutive  
order for A, B, C

$$MT(B) = O(1) \leftarrow \text{base}$$

$$MT(c\sqrt{M}) = O\left(\frac{M}{B}\right)$$



$$\Rightarrow MT(N) = O\left(\frac{N^3}{M^{3/2}} \cdot \frac{M}{B}\right) = O\left(\frac{N^3}{B\sqrt{M}}\right) \quad \text{Optimal!}$$

in the two-level  
memory model.